

Problem:

# FRESHLLMs: REFRESHING LARGE LANGUAGE MODELS WITH SEARCH ENGINE AUGMENTATION side from!

Tu Vu<sup>1</sup> Mohit Iyyer<sup>2</sup> Xuezhi Wang<sup>1</sup> Noah Constant<sup>1</sup> Jerry Wei<sup>1</sup>  
 Jason Wei<sup>3\*</sup> Chris Tar<sup>1</sup> Yun-Hsuan Sung<sup>1</sup> Denny Zhou<sup>1</sup> Quoc Le<sup>1</sup> Thang Luong<sup>1</sup>  
 Google<sup>1</sup> University of Massachusetts Amherst<sup>2</sup> OpenAI<sup>3</sup> How to use:  
 freshllms@google.com

## ABSTRACT

Most large language models (LLMs) are trained once and never updated; thus, they lack the ability to dynamically adapt to our ever-changing world. In this work, we perform a detailed study of the factuality of LLM-generated text in the context of answering questions that test current world knowledge. Specifically, we introduce FRESHQA, a novel dynamic QA benchmark encompassing a diverse range of question and answer types, including questions that require *fast-changing* world knowledge as well as questions with *false premises* that need to be debunked. We benchmark a diverse array of both closed and open-source LLMs under a two-mode evaluation procedure that allows us to measure both correctness and hallucination. Through human evaluations involving more than 50K judgments, we shed light on limitations of these models and demonstrate significant room for improvement: for instance, all models (regardless of model size) struggle on questions that involve fast-changing knowledge and false premises. Motivated by these results, we present FRESHPROMPT, a simple few-shot prompting method that substantially boosts the performance of an LLM on FRESHQA by incorporating relevant and up-to-date information retrieved from a search engine into the prompt. Our experiments show that FRESHPROMPT outperforms both competing search engine-augmented prompting methods such as SELF-ASK (Press et al., 2022) as well as commercial systems such as PERPLEXITY.AI.<sup>1</sup> Further analysis of FRESHPROMPT reveals that both the number of retrieved evidences and their order play a key role in influencing the correctness of LLM-generated answers. Additionally, instructing the LLM to generate concise and direct answers helps reduce hallucination compared to encouraging more verbose answers. To facilitate future work, we release FRESHQA at [github.com/freshllms/freshqa](https://github.com/freshllms/freshqa) and commit to updating it at regular intervals.

## 1 INTRODUCTION

Recent large language models (LLMs) such as BARD and CHATGPT/GPT-4<sup>2</sup> are designed to be versatile open-domain *chatbots* that can engage in multi-turn conversations on diverse subjects. Despite their impressive capabilities, these LLMs often “hallucinate” plausible but factually incorrect information (Maynez et al., 2020; Liu et al., 2023b), which reduces the trustworthiness of their responses, especially in settings where accurate and up-to-date information is critical. This behavior can be partially attributed to the presence of outdated knowledge encoded in their parameters. While additional training using human feedback (Ouyang et al., 2022) or knowledge-enhanced tasks can mitigate this issue, it is not easily scalable for real-time knowledge updates (e.g., stock price of a company). In-context learning (Brown et al., 2020) is an appealing alternative in which real-time knowledge can be injected into an LLM’s prompt for conditioning generation. While recent work has begun to explore augmenting LLMs with web search results (Lazaridou et al., 2022; Press et al., 2022), it is unclear how to take full advantage of search engine outputs to increase LLM factuality.

\*Work done while at Google.

<sup>1</sup><https://www.perplexity.ai>

<sup>2</sup><https://bard.google.com>, <https://chat.openai.com>

Type	Question	Answer (as of this writing)
never-changing	Has Virginia Woolf's novel about the Ramsay family entered the public domain in the United States?	<b>Yes</b> , Virginia Woolf's 1927 novel <i>To the Lighthouse</i> entered the public domain in 2023.
never-changing	What breed of dog was Queen Elizabeth II of England famous for keeping?	<b>Pembroke Welsh Corgi</b> dogs.
slow-changing	How many car models does Tesla offer?	Tesla offers <b>four</b> car models: Model S, Model X, Model 3 and Model Y.
slow-changing	Which team holds the record for largest deficit overcome to win an NFL game?	The record for the largest NFL comeback is held by the <b>Minnesota Vikings</b> .
fast-changing	Which game won the Spiel des Jahres award most recently?	<b>Cascadia</b> won the 2022 Spiel des Jahres.
fast-changing	What is Brad Pitt's most recent movie as an actor	Brad Pitt recently starred in <b>Babylon</b> , directed by Damien Chazelle.
false-premise	What did Donald Trump's first Tweet say after he was unbanned from Twitter by Elon Musk?	He <b>has not yet tweeted</b> since he was unbanned.
false-premise	In which round did Novak Djokovic lose at the 2022 Australian Open?	He <b>was not allowed to play</b> at the tournament due to his vaccination status.

Figure 1: FRESHQA exemplars. Our questions are broadly divided into *four* main categories based on the nature of the answer: *never-changing*, in which the answer almost never changes; *slow-changing*, in which the answer typically changes over the course of several years; *fast-changing*, in which the answer typically changes within a year or less; and *false-premise*, which includes questions whose premises are factually incorrect and thus have to be rebutted.

In this work, we collect a novel QA benchmark, dubbed FRESHQA, to evaluate the factuality of existing LLMs. FRESHQA consists of 600 natural questions that are broadly divided into the *four* main categories shown in Figure 1. FRESHQA’s questions span a diverse set of topics with diverse difficulty levels (requiring single-hop and multi-hop reasoning), and require a model to “understand” the world’s up-to-date knowledge to be able to answer correctly. Additionally, FRESHQA is dynamic in nature: some of the ground-truth answers may change over time, and a question classified under a specific category may undergo reclassification at some later point in time (e.g., the current *false-premise* question “How long has Elon Musk been married to his current spouse?” will fall into the *fast-changing* category if Elon Musk gets married again in the future).

We benchmark how well different LLMs perform on FRESHQA by prompting them with questions and optionally a few question-answer demonstrations and then sampling a response. Then, we conduct an extensive human evaluation of the factual accuracy of the models’ responses, consisting of more than 50K judgements. We evaluate each response in a two-mode evaluation procedure: RELAXED, which measures only whether the main answer is correct; and STRICT, which measures whether all of the claims in the response are factual and up-to-date (i.e., no hallucination). Our study sheds light on the factuality of old and new LLMs and reveals different model behaviors across question types. Unsurprisingly, there are flat scaling curves on questions that involve fast-changing knowledge: simply increasing the model size does not lead to reliable performance gains. We also observe similar trends on false-premise questions, though several LLMs are able to debunk a false-premise question if explicitly asked “Please check if the question contains a valid premise before answering”. Overall, FRESHQA is challenging for current LLMs and leaves ample room for improvement.

Motivated by these findings, we further investigate how to effectively improve LLMs’ factuality by grounding their responses to accurate and up-to-date information from search engines. Given the rapid development of ever larger LLMs and the ever-changing nature of knowledge, we explore in-context learning approaches that allow an LLM to attend over knowledge provided at inference time through its prompt. We develop FRESHPROMPT, a simple yet effective method that, for a given question, takes full advantage of a search engine by extracting all up-to-date and relevant information (including knowledge from relevant questions that search users also ask) and uses few-shot in-context learning to teach a model to reason over retrieved evidences and figure out the right answer. We show that FRESHPROMPT significantly boosts LLMs’s factuality: for example, our best GPT-4 + FRESHPROMPT variant yields an improvement of 32.6% and 49.0% accuracy over the vanilla GPT-4 on FRESHQA under RELAXED and STRICT, respectively. Since our method requires no additional training, it is flexible and applicable to a variety of scenarios.

Taken together, our key contributions include:

- We introduce a novel dynamic QA benchmark, *FRESHQA*, which features a diverse set of question and answer types, including questions whose answers may change over time and questions whose premises are factually incorrect. We will make our dataset freely available and commit to updating the ground-truth answers at a regular schedule to encourage exploration of methods to improve LLMs’ factuality.
- We benchmark a wide range of both closed and open-source LLMs on our dataset. Through an extensive and rigorous human evaluation study, we shed light on limitations of current LLMs: they struggle on fast-changing, false-premise, and multi-hop questions, and our two-mode evaluation captures increased hallucinations produced by techniques such as chain-of-thought prompting (Wei et al., 2022).
- We present *FRESHPROMPT*, a simple in-context learning method that can substantially boost an LLM’s factuality compared to competing search-augmented approaches by effectively incorporating factual and up-to-date information from a search engine into the model’s prompt. Furthermore, we perform a series of sensitivity and ablation analyses to better understand what facets of *FRESHPROMPT* contribute to its success.

## 2 FRESHQA

In this section, we address the growing need to assess LLM factuality by curating a novel QA benchmark, *FRESHQA*, with 600 questions that cover a wide spectrum of question and answer types.

### 2.1 DATA COLLECTION

We collected *FRESHQA* by recruiting both NLP researchers (including the authors and their colleagues) and online freelancers<sup>3</sup> to write questions of varying difficulty levels and topics whose answers may change based on new developments in the world. The annotators were shown a few exemplars of the four broad types of questions defined in Figure 1. Within each of these four categories, we ask annotators to write questions at two different difficulty levels: *one-hop*, where the question explicitly mentions all of the relevant information needed to answer it, and thus no additional reasoning is required (e.g., “*Who is the CEO of Twitter?*”); and *multi-hop*, where the question requires one or more additional steps of reasoning in order to gather all of the relevant information needed to answer it (e.g., “*What is the total height of the tallest building in the world?*”). Annotators were encouraged to write questions that involve *fresh* knowledge (knowledge that has changed recently or new events) and appear *natural* (i.e., plausible for a real person to type into a search engine). For false-premise questions, we requested a brief explanation elucidating why the question is flawed.<sup>4</sup>

**Quality control:** Upon obtaining the initial dataset, we conducted multiple thorough data cleaning and quality assessments. This involved manual review of each example to ensure well-formed questions, removal of duplicates and invalid questions (e.g., too easy or controversial), and verification of answers and supporting evidence URLs. We also manually collected supplementary valid answers for each question (e.g., different names of the same person, different date formats, etc.). To facilitate future answer updates, we excluded questions whose answers are likely to change more frequently than once per week, and additionally incorporated the expected next review date for each question.

**Data size and split:** The resulting dataset is divided into a *test* set consisting of 125 questions for each of the four broad question types (500 total examples) and a *development* set comprising 25 questions for each question type (100 total examples), sampled randomly within types. Additionally, 15 examples spanning different question types were extracted for *demonstration* purposes (i.e., for use in few-shot in-context learning), and the remaining data was discarded. The development set is reserved for future studies and not used in this paper.<sup>5</sup>

**FRESHQA requires regular updates:** Our dataset has time sensitivity since the ground-truth answers may change with new developments in the world. As such, we commit to updating the dataset regularly and encourage researchers to evaluate on the latest version of the dataset, as close to the release date of the updated dataset as possible.

<sup>3</sup>We use Upwork (<https://www.upwork.com>) with a compensation rate of \$2 per example.

<sup>4</sup>Additionally, the annotators were asked to include the year the answer to the question last changed and a URL to a reputable website that supports the answer.

<sup>5</sup>Although our test set is currently balanced across question types, the distribution may change over time due to reclassification of questions from one category to another.

## 2.2 EVALUATION

All model responses were evaluated by the authors in a two-mode evaluation procedure: **RELAXED**, which focuses solely on evaluating the correctness of the primary answer; and **STRICT**, which additionally examines whether *all* of the facts in the answer are accurate (i.e., no hallucination). Overall, our setup provides both ends of the spectrum for evaluating factuality (the difference between a model’s strict and relaxed performance provides a way to measure hallucination), offering a more comprehensive and nuanced understanding of their performance. Two authors independently evaluated a subset of 100 answers in both modes and had an agreement of 99% for **RELAXED** and 96% for **STRICT**, showing that the protocol is reliable for comparing different LLMs (see Appendix A).

**Evaluation protocol:** In both evaluation modes, we credit a model’s response only if it provides a confident and definitive answer, or the correct answer can be obviously inferred from the response. The primary or final answer when standing alone must be accurate. Any additional information that is provided must not contradict the primary answer or reshape one’s perception of it. For false-premise questions, the model must point out the presence of a false premise to receive credit. For answers that involve names of entities (e.g., people), complete names or commonly recognized names are expected. Regarding numerical answers, approximate numbers are generally not accepted unless explicitly included in the ground-truth answers. In **RELAXED**, we accept ill-formed responses (including those in a non-English language), as well as hallucinated or outdated information that does not significantly impact the primary answer. Under **STRICT**, however, a response that contains any hallucination, no matter how minor, will not receive credit. Furthermore, we accept a response in **STRICT** when the model indicates that the information might be outdated (e.g., “*As of my knowledge cutoff date in September 2021*”) *only* if it is evident that the knowledge has not changed.<sup>6</sup> Figure 4 in Appendix B shows specific examples of each evaluation criteria.

## 3 PRE-TRAINED LLMs STRUGGLE ON FRESHQA

We use **FRESHQA** to benchmark LLMs that do not have access to real-time data or the ability to browse the Internet for current information.<sup>7</sup> While all LLMs (regardless of size) predictably struggle on questions requiring up-to-date knowledge, they also underperform on false premise questions. In our experiments, we simply feed individual questions as prompts into each model and decode the model’s predictions using a temperature of 0 without fine-tuning (see Appendix C for more details).

**Baselines:** We experiment with a series of models varying in size from 770M to 540B parameters, including basic pre-trained models such as **T5** (Raffel et al., 2020; Lester et al., 2021), **PALM** and **PALMCHILLA** (Chowdhery et al., 2022), optionally using **FEW-SHOT** prompting (Brown et al., 2020) and Chain-of-Thought (**COT**, Wei et al., 2022),<sup>8</sup> instruction-tuned models including **FLAN-T5** and **FLAN-PALM** (Chung et al., 2022; Longpre et al., 2023), and OpenAI’s **GPT-3.5** (Ouyang et al., 2022), **CODEX** (Chen et al., 2021a), **CHATGPT**, and **GPT-4** (OpenAI, 2023).

### 3.1 RESULTS AND DISCUSSION

**FRESHQA presents a challenge for LLMs:** We visualize the accuracy of different LLMs on **FRESHQA** in both evaluation modes in Figure 2.<sup>9</sup> A first obvious takeaway is that all models struggle on **FRESHQA**: overall accuracy ranges from 0.8% to 32.0% under **STRICT**, and 0.8% to 46.4% under **RELAXED**. Switching from **RELAXED** to **STRICT** results in a marked decrease in accuracy for **CHATGPT** and **GPT-4**. This is mainly due to the lack of access to up-to-date information, as they produce “outdated” answers (which often start with the prefix “*As of my knowledge cutoff date in*

<sup>6</sup>Note that even without access to real-time data, a model may still provide accurate answers to certain questions involving current information, potentially through random guesses or by leveraging past valid responses (e.g., for the question “Which drama series won the most recent Primetime Emmy Award for Outstanding Drama Series?”, while “*Succession*” won the award most recently (as of this writing), it was also the winner in 2020, so a model trained in 2021 could potentially provide the correct answer.).

<sup>7</sup>With the exception of **CHATGPT** and **GPT-4**, which have access to the current date. Note that the latest versions of these models can now browse the Internet.

<sup>8</sup>As we are interested in exploring how these methods perform without being specifically designed for **FRESHQA**, we use the 5-shot demonstrations for **TRIVIAQA** (Joshi et al., 2017) used in Sun et al. (2023).

<sup>9</sup>Table 3 and Table 4 in Appendix D contain concrete numbers under **STRICT** and **RELAXED**, respectively.




Figure 2: Accuracy of different LLMs on FRESHQA under RELAXED and STRICT (no hallucination) evaluations. **Models benchmarked on the same date of April 26, 2023.** All models (regardless of model size) struggle on questions that involve *fast-changing* knowledge and *false premises*.

September 2021”), and in many cases, “refuse” to provide an answer (e.g., “As an AI language model, I cannot provide real-time information.”). Similarly, the accuracy of PALM (across model sizes) drops significantly under STRICT. Much of this drop is due to artifacts such as conversation-like responses with unexpected special tokens (e.g., the end-of-turn [eot]), and hallucination. In contrast, FLAN-PALM and CODEX exhibit minimal hallucination due to their concise and direct answers.

**LLMs struggle with questions about current information:** The lack of up-to-date parametric knowledge results in dramatically degraded accuracies across models on questions involving fast-changing or recent knowledge. GPT-4 generally obtains the highest accuracy on these questions, with the exception of questions about recent knowledge (i.e., since 2022) under STRICT where it underperforms FLAN-PALM and CODEX, but it never exceeds 15% across both evaluation modes. Our evaluation confirms that CHATGPT and GPT-4 have been exposed to data containing information beyond their knowledge cutoff date (Appendix E). Additionally, GPT-4 is more reluctant to answer fast-changing questions (refusing to answer 60% of the time) compared to CHATGPT (16%).

**Questions with false premises pose a hurdle for LLMs:** All models struggle on questions with false premises, and using larger models does not increase accuracy for T5 and PALM (“flat scaling”), with performance within the range of 0.0% to 1.6%. GPT-3.5, CHATGPT, and GPT-4 demonstrate much superior accuracies to all other models, achieving accuracies between 25.8% to 42.7% under STRICT and 32.3% to 66.9% under RELAXED. CHATGPT performs the best under STRICT (42.7%) while GPT-4 is the most accurate model under RELAXED (66.9%), with an impressive accuracy of 83.9% on questions about knowledge before 2022. These results suggest that OpenAI’s models are likely trained to cope with false-premise questions.

**CoT increases hallucination:** Overall, FEW-SHOT and CoT prompting are beneficial for large models and sometimes advantageous for moderately-sized models on questions with valid premises, especially on questions about never-changing or old knowledge. Under STRICT, FEW-SHOT and CoT yields +36.1% and +26.9% respective accuracy improvement over zero-shot prompting with PALM 540B on questions involving knowledge before 2022 (+21.9% and +29.7% under RELAXED). CoT largely demonstrates superior performance compared to FEW-SHOT under RELAXED, whereas FEW-SHOT obtains better results under STRICT, as CoT introduces more room for hallucination.

**Multi-hop reasoning is challenging for several models:** T5 LARGE and XL are incapable of dealing with multi-hop questions, while FLAN-PALM 540B, CODEX, and GPT-3.5 suffer the most when switching from one-hop to multi-hop questions. GPT-4 remains stable across these two types of questions (with a difference of less than 2% in accuracy across settings). See Appendix D for details.

## 4 PROMPTING SEARCH ENGINE-AUGMENTED LANGUAGE MODELS

The low accuracies reported in the previous section are largely unsurprising, as none of the models we evaluated had access to real-time information. In this section, we evaluate the impact of *search engine augmentation* to LLMs on FRESHQA. We present FRESHPROMPT, a simple few-shot prompting method that substantially boosts FRESHQA performance of an LLM by incorporating relevant and up-to-date information retrieved from a search engine (GOOGLE SEARCH) into the prompt.

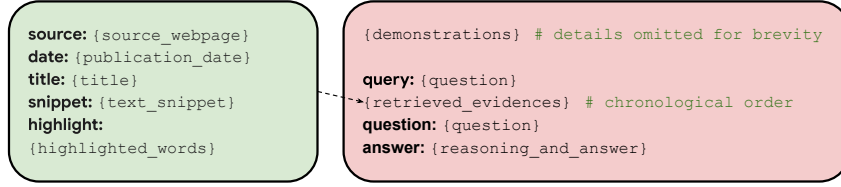


Figure 3: FRESHPROMPT’s format. We cast all retrieved evidences into a unified format with useful information, including source webpage, date, title, text snippet, and highlighted words (left). Few-shot demonstrations are provided at the beginning of the prompt. Each demonstration shows the model an example question and a list of retrieved evidences for the question, followed by some reasoning over the evidences to figure out the most relevant and up-to-date answer (right).

#### 4.1 FRESHPROMPT

Our FRESHPROMPT method leverages a text prompt to (1) introduce contextually relevant and up-to-date information (including answers to relevant questions) from a search engine to a pre-trained LLM, and (2) teach the model to reason over retrieved evidences. More specifically, given a *question*  $q$ , we first use  $q$  verbatim to query a search engine, in our case GOOGLE SEARCH.<sup>10</sup> We retrieve all of the search results, including the *answer box*, *organic results*, and other useful information, such as the *knowledge graph*, *questions and answers* from crowdsourced QA platforms, and *related questions* that search users also ask (see Figure 6 in Appendix F). For each of these results, we extract the associated *text snippet*  $x$  along with other information, such as *source*  $s$  (e.g., WIKIPEDIA), *date*  $d$ , *title*  $t$ , *highlighted words*  $h$ , and then create a list of  $k$  retrieved evidences  $E = \{(s, d, t, x, h)\}$ . These evidences are then cast into a common format (Figure 3, left) and used to condition the model through in-context learning. To encourage the model to focus on more recent evidences following recent findings (Liu et al., 2023a), we sort the evidences  $E$  in the prompt from oldest to newest.

To help the model to “understand” the task and the desired output, we provide few-shot demonstrations of input-output exemplars at the beginning of the input prompt. Each demonstration shows the model an example question and a list of retrieved evidences for the question, followed by a chain-of-thought reasoning over the evidences to figure out the most relevant and up-to-date answer (Figure 3, right). Although we include a few exemplars of questions with false premises in the demonstrations, we also experiment with an explicit false premise check in the prompt: “Please check if the question contains a valid premise before answering”. Figure 7 in Appendix G shows a realistic prompt.

#### 4.2 EXPERIMENT SETUP

We closely follow the setup in Section 3 except in cases where we lack control over the model’s decoding via an API (e.g., PERPLEXITY AI). Some of the models we evaluate can potentially change over time, which presents a challenge to the reproducibility of our evaluation results; thus, we evaluate all models on the same date of April 26, 2023. In addition to GPT-3.5 and GPT-4, we evaluate GOOGLE SEARCH by simply querying GOOGLE SEARCH and using the answer in the answer box (if any) or the text snippet of the top-1 search result; PERPLEXITY.AI (PPL.AI), an answer engine that combines a LLM and a search engine to generate useful responses to users’ queries;<sup>11</sup> and SELF-ASK (Press et al., 2022), a method that uses few-shot in-context learning to teach a LLM to decompose each question into simpler sub-questions that are answered via GOOGLE SEARCH.<sup>12</sup>

**FRESHPROMPT setup:** We apply FRESHPROMPT to both GPT-3.5 and GPT-4 by sequentially incorporating the following retrieved evidences into the input prompt:  $o$  organic search results,  $r$  related questions that search users also ask,  $a$  questions and answers from crowdsourced QA platforms, and the snippets from the knowledge graph and answer box (if available). These evidences are arranged in sequence up to the end of the prompt. Given the models’ context limit, we only keep the top  $n$  evidences (closer to the end of the prompt) after sorting them based on the cor-

<sup>10</sup>We scrape the results from GOOGLE SEARCH using SERPAPI (<https://serpapi.com>).

<sup>11</sup><https://www.perplexity.ai>. At the time of evaluation, PPLX.AI was a combination of GPT-3.5 and BING SEARCH, and was able to provide both concise and detailed answers. We evaluated its concise answers.

<sup>12</sup>We use the few-shot prompt provided by SELF-ASK’s authors and apply it to both GPT-3.5 and GPT-4. For simplicity, we evaluate solely the final answer from SELF-ASK, disregarding intermediate answers.



Table 1: Accuracy of different search engine-augmented LLMs on FRESHQA under STRICT (no hallucination) evaluations. **Models benchmarked on the same date of April 26, 2023.** We report accuracy across different categories of questions, including *fast-changing* (*fast*), *slow-changing* (*slow*), *never-changing* (*never*), false-premise, questions that involve knowledge before 2022 ( $< 2022$ ) and since 2022 ( $\geq 2022$ ), one-hop (*1-hop*) and multi-hop (*m-hop*) questions. <sup>+</sup> indicates a model with access to the current date. UTD stands for “up-to-date”.

Model (size)	knowl. cutoff	all	valid premise							false premise		
			all	fast	slow	never	< 2022	≥ 2022	1-hop	m-hop	all	< 2022
<i>comparison against baselines</i>												
GOOGLE SEARCH (N/A)	UTD	39.6	48.9	32.0	46.4	68.3	67.4	37.9	55.6	32.4	11.3	9.7
GPT-3.5 (N/A)	2021	26.0	26.1	4.0	15.2	58.7	61.0	5.1	28.0	21.3	25.8	34.4
GPT-3.5 + SELF-Ask (N/A)	UTD	41.6	51.1	36.8	43.2	73.0	73.8	37.4	52.2	48.1	12.9	17.2
GPT-3.5 + FRESHPROMPT	UTD	56.0	62.5	46.4	60.8	80.2	71.6	57.0	68.7	47.2	36.3	43.0
PPLX.AI (N/A)	UTD	52.2	57.2	38.4	53.6	79.4	73.0	47.7	63.8	40.7	37.1	38.7
GPT-4 (N/A)	2021 <sup>+</sup>	28.6	26.9	12.0	4.0	64.3	58.2	8.1	27.2	25.9	33.9	41.9
GPT-4 + SELF-Ask (N/A)	UTD	47.8	47.1	39.2	46.4	55.6	51.8	44.3	43.7	55.6	50.0	61.3
GPT-4 + FRESHPROMPT	UTD	<b>75.6</b>	<b>77.1</b>	<b>59.2</b>	<b>77.6</b>	<b>94.4</b>	<b>88.7</b>	<b>70.2</b>	<b>81.3</b>	<b>66.7</b>	<b>71.0</b>	<b>77.4</b>
<i>sensitivity and ablation studies</i>												
GPT-3.5 (N/A)	2021	26.0	26.1	4.0	15.2	58.7	61.0	5.1	28.0	21.3	25.8	34.4
GPT-3.5 + FRESHPROMPT	UTD	56.0	62.5	46.4	60.8	80.2	71.6	57.0	68.7	47.2	36.3	43.0
w/ PREMISE CHECK	UTD	35.2	27.1	14.4	28.0	38.9	36.2	21.7	31.0	17.6	59.7	67.7
GPT-4 (N/A)	2021 <sup>+</sup>	28.6	26.9	12.0	4.0	64.3	58.2	8.1	27.2	25.9	33.9	41.9
GPT-4 w/ SNIPPETS ONLY & SEARCH ORDER	UTD	74.0	75.5	56.8	75.2	94.4	87.9	68.1	79.9	64.8	69.4	77.4
GPT-4 w/ SNIPPETS ONLY & TIME ORDER	UTD	74.8	75.5	58.4	74.4	93.7	87.9	68.1	79.9	64.8	72.6	<b>82.8</b>
GPT-4 w/ SNIPPETS ONLY & RANDOM ORDER	UTD	72.4	73.7	56.8	69.6	94.4	87.9	65.1	78.4	62.0	68.5	76.3
GPT-4 + FRESHPROMPT	UTD	75.6	77.1	59.2	77.6	94.4	<b>88.7</b>	70.2	81.3	66.7	71.0	77.4
w/ PREMISE CHECK	UTD	75.0	74.2	56.8	76.0	89.7	85.1	67.7	79.5	61.1	<b>77.4</b>	79.6
w/o ANSWER BOX	UTD	74.2	74.7	57.6	74.4	92.1	<b>88.7</b>	66.4	79.1	63.9	72.6	78.5
w/o ANSWER BOX & RELEVANT INFO	UTD	72.4	72.9	54.4	71.2	92.9	87.2	64.3	78.0	60.2	71.0	78.5
w/ 1 EVIDENCE	UTD	61.4	60.9	40.0	55.2	87.3	79.4	49.8	66.8	46.3	62.9	75.3
w/ 5 EVIDENCES	UTD	70.6	72.1	56.0	69.6	90.5	81.6	66.4	78.0	57.4	66.1	73.1
w/ 15 EVIDENCES	UTD	<b>77.6</b>	<b>78.5</b>	<b>60.8</b>	<b>78.4</b>	<b>96.0</b>	<b>88.7</b>	<b>72.3</b>	<b>81.7</b>	<b>70.4</b>	75.0	80.6
w/ 15 DEMONSTRATIONS	UTD	74.6	75.5	56.8	76.0	93.7	87.9	68.1	79.9	64.8	71.8	76.3
w/ LONG DEMONSTRATION ANSWERS	UTD	73.0	72.6	55.2	71.2	91.3	83.7	66.0	77.6	60.2	74.2	81.7

responding date. Unless otherwise specified, we use  $(o, r, a, n, m) = (10, 2, 2, 5)$  for GPT-3.5, and  $(o, r, a, n, m) = (10, 3, 3, 10)$  for GPT-4. Additionally, we include  $m = 5$  question-answer demonstrations at the beginning of the prompt.

### 4.3 RESULTS AND DISCUSSION

**FRESHPROMPT significantly improves FRESHQA accuracy:** Table 1 presents concrete numbers under STRICT (see Appendix H for results under RELAXED). FRESHPROMPT offers large improvements over the vanilla GPT-3.5 and GPT-4 across the board. GPT-4 + FRESHPROMPT achieves absolute accuracy improvements of 47% and 31.4% over GPT-4 under STRICT and RELAXED, respectively. The reduction in the absolute accuracy gap between STRICT and RELAXED (from 17.8% to 2.2%) also suggests that FRESHPROMPT dramatically diminishes the presence of outdated and hallucinated answers. Unsurprisingly, the most significant improvements for both GPT-3.5 and GPT-4 are on the categories of fast-changing and slow-changing questions, which both concern recent knowledge. That said, questions about old knowledge also benefit from FRESHPROMPT. For example, GPT-4 + FRESHPROMPT yields a +30.5% higher accuracy than GPT-4 on questions with valid premises that involve knowledge before 2022 (+9.9% under RELAXED). Additionally, FRESHPROMPT produces notable gains on false-premise questions (+37.1% and +8.1% respective accuracy improvements under STRICT and RELAXED for GPT-4).

**FRESHPROMPT outperforms other search-augmented methods by a large margin:** GPT-4 + FRESHPROMPT demonstrates superior accuracy across question types, surpassing all other methods by a substantial margin. Its best variant (with 15 retrieved evidences per question) achieves impressive overall accuracies of 77.6% and 79.0% under STRICT and RELAXED, respectively. GPT-3.5 + FRESHPROMPT surpasses PPLX.AI and SELF-ASK (all performed on top of GPT-3.5) in overall accuracy by +3.8% and +14.4% under STRICT. Under RELAXED, however, PPLX.AI achieves a +4.2% higher overall accuracy than GPT-3.5 + FRESHPROMPT, which is a large part due to its superior accuracy on false-premise questions (58.1% vs. 41.1%). The large accuracy gap of 14.0% between STRICT and RELAXED for PPLX.AI show that the method contains a large amount of hallucination. Overall, all search-engine augmented approaches (SELF-ASK, PPLX.AI, and FRESHPROMPT) provide significant

gains across question types over vanilla GPT-3.5 and GPT-4. GOOGLE SEARCH generally provides better results than both GPT-3.5 and GPT-4, except on questions with false premises, but lags far behind PPLX.AI and GPT-3.5/GPT-4 + FRESHPROMPT across the board.

**The premise check boosts accuracy on false-premise questions but can hurt accuracy on those with valid premises:** As discussed in Section 3.1, OpenAI’s LLMs such as GPT-3.5 and GPT-4 are likely tuned to handle false-premise questions, and this is also true for PPLX.AI. Additionally, we empirically find that several LLMs possess the ability to debunk a false-premise question if explicitly asked, e.g., “*Please check if the question contains a valid premise before answering*”. Adding this premise check to GPT-3.5 and GPT-4 yields +23.4% and +6.4% respective accuracy improvement on false-premise questions under STRICT (+22.6% and +11.3% under RELAXED). However, this is harmful for GPT-3.5 with regard to other question types, decreasing overall accuracy by 20.8% and 21% under STRICT and RELAXED, respectively. This is not a problem for GPT-4, with a slight decrease of 0.6% under STRICT and a slight increase of 1.2% under RELAXED.

**Having more relevant and up-to-date evidences at the end of the input context is helpful:** We also analyze how the order of the evidences in the prompt impacts GPT-4’s accuracy. Our results show that using the order returned by GOOGLE SEARCH (SEARCH ORDER, top search results at the end of the input context) or sorting the evidences by their associated date information (TIME ORDER, more recent results at the end) always results in slightly better accuracy compared to using a random order (RANDOM ORDER), with up to a +2.2% higher overall accuracy in STRICT and RELAXED. Using only the text snippet for each evidence without additional information (such as source, date, etc.) as in GPT-4 + FRESHPROMPT slightly reduces accuracy, with less than 1% in both settings.

**Additional retrieved information beyond the organic search results provides further gains:** Incorporating additional retrieved evidences other than the *organic search results*, such as the *answer box* or *related questions* that search users also ask, is helpful. Removing the *answer box* decreases GPT-4 + FRESHPROMPT’s overall accuracy under STRICT by 1.4% (1.6% under RELAXED). Removing both the *answer box* and other relevant information (including *related questions*) reduces GPT-4 + FRESHPROMPT’s overall accuracy by 3.2% (3.0% under RELAXED).

**Increasing the number of retrieved evidences further improves FRESHPROMPT:** We explore the effect of the number of retrieved evidences for each question as well as the number of demonstrations by varying these numbers in our experiments with GPT-4. Note that our default setting for GPT-4 + FRESHPROMPT uses 10 retrieved evidences for each question and 5 demonstrations. Our results suggest that the number of retrieved evidences for each question is the most important ingredient for achieving highest accuracy. Under STRICT, increasing this number from 1 to 5, 10, and 15 leads to corresponding overall accuracy improvements of +9.2%, +14.2%, and +16.2%, respectively. This suggests that GPT-4 is able to efficiently handle an increasing number of retrieved evidences (including conflicting answers) and ground its responses into the most factual and up-to-date information. On the other hand, increasing the number of *demonstrations* from 5 to 15 slightly hurts accuracy in both evaluation settings (1% decrease in overall accuracy under STRICT).

**Verbose demonstrations improve on complex questions but also increase hallucination:** To evaluate the effect of the writing style of the answer (including the reasoning) in each demonstration, we manually rewrite these answers into a more verbose version (LONG DEMONSTRATION ANSWERS). Our manual inspection reveals that using more verbose demonstration answers may be helpful when dealing with complex questions but can be more harmful as it provides room for hallucination (a decrease of 2.6% in overall accuracy under STRICT).

## 5 RELATED WORK

**Knowledge augmented LLMs:** Many prior works study semi-parametric knowledge augmentation in LLMs via additional fine-tuning (Guu et al., 2020; Lewis et al., 2020; Borgeaud et al., 2022; Izacard et al., 2022), while others advocate for knowledge generation instead of retrieval (Yu et al., 2023a; Sun et al., 2023). FRESHPROMPT aligns with a recent emerging trend in QA applications that augments LLMs’ prompts with knowledge retrieved from search engines for real-time alignment to current and factual information (Nakano et al., 2021; Lazaridou et al., 2022; Menick et al., 2022; Yao et al., 2022; Press et al., 2022; Khattab et al., 2022; Schick et al., 2023; Luo et al., 2023). Similar to our method, Lazaridou et al. (2022) proposed a few-shot in-context learning approach that inserts documents from GOOGLE SEARCH into the prompt. We do not compare to this method due to its expensive inference cost, as it chunks retrieved documents into evidence paragraphs and performs



$k = 50$  inference calls to the LLM to generate  $k$  answers followed by LLM reranking. In contrast, FRESHPROMPT only performs a single inference call to the LLM. SELF-ASK (Press et al., 2022) also uses few-shot in-context learning to teach an LLM to ask itself follow-up questions before answering the initial question, although it focuses more on decomposition.

**Time-sensitive QA:** FRESHQA aligns with a growing body of work on benchmarking LLMs’ temporal reasoning capabilities (Chen et al., 2021b; Zhang & Choi, 2021; Liska et al., 2022; Kasai et al., 2022). Chen et al. (2021b) created TIMEQA by extracting evolving facts from WIKIDATA along with aligned WIKIPEDIA passages to synthesize 20K timestamped question-answer pairs. Zhang & Choi (2021) constructed SITUATEDQA by annotating 9K realistic questions from existing open-domain QA datasets with temporal context (i.e., timestamps). STREAMINGQA (Liska et al., 2022) consists of both LLM-generated and human-written questions (146K total questions) answerable from a corpus of timestamped news articles. Also related is the dynamic REALTIMEQA benchmark (Kasai et al., 2022), which evaluates models on a set of 30 multiple-choice questions about new events extracted from news websites. In contrast, FRESHQA contains a fixed set of human written open-ended questions whose answers by nature can change based on new developments in the world and thus offers a complementary generative evaluation of time-sensitive QA.

**QA over questionable or counterfactual premises:** Recent work has also introduced QA benchmarks with questionable premises (Yu et al., 2023c; Kim et al., 2023) or counterfactual premises (Yu et al., 2023b). CREPE (Yu et al., 2023c) consists of 8400 Reddit questions (of which 25% questions contain false premises annotated by human workers) split into train/dev/test sets. Kim et al. (2023) constructed (QA)<sup>2</sup>, an evaluation set of 602 questions based on frequent search engine queries, which are annotated by expert annotators and crowdworkers, and evenly divided between those with and without questionable premises. Consistent with these efforts, we find that current LLMs struggle with handling false premise questions; additionally, several LLMs are able to debunk a false-premise question if explicitly asked to check for the premise’s validity. Similar to above, these benchmarks are complementary and combining them is a promising direction for future work.

## 6 LIMITATIONS AND FUTURE WORK

One obvious challenge with FRESHQA is the need for regular answer updating by the maintainers; in the interim period between updates, the answers to some questions might become stale. This could be addressed by support from the open-source community (e.g., updates via GITHUB pull requests). Also, STRICT human evaluation is currently expensive, as it requires the verification of all claims in each generated response; this could be addressed by implementing automatic LLM-based evaluators. On the method side, FRESHPROMPT interfaces with GOOGLE SEARCH, and it is unclear how it performs with other search engines for which some types of context (e.g., answer boxes) are not available. Additionally, we only perform one search query per question, and thus our method could be further improved via question decomposition and multiple search queries (Khattab et al., 2022). Since FRESHQA consists of relatively simple English language questions, it is also unclear how well FRESHPROMPT performs in the context of multilingual/cross-lingual QA and long-form QA (Fan et al., 2019). Finally, FRESHPROMPT relies on in-context learning and thus may underperform approaches that fine-tune the base LLM on new knowledge.

## 7 CONCLUSION

Our work offers a fine-grained and exhaustive evaluation of the capabilities of modern LLMs to adapt to ever-changing world knowledge with and without search engine augmentation. In the process, we develop a new dataset—FRESHQA—of 600 questions that test a broad range of reasoning abilities, from the incorporation of fast-changing knowledge to identification of questions with false premises. Our two-mode evaluation also provides a way to measure both correctness and hallucination. Additionally, we propose a simple few-shot in-context learning algorithm called FRESHPROMPT that incorporates relevant evidences retrieved from GOOGLE SEARCH into the prompt of an LLM. FRESHPROMPT significantly improves performance over competing search engine-augmented approaches on FRESHQA, and an ablation reveals that factors such as the number of incorporated evidences and their order impact the correctness of LLM-generated answers. We release FRESHQA and commit to updating its answers regularly to facilitate future research.

## 8 ACKNOWLEDGEMENTS

We thank Colin Raffel, Hamed Zamani, and Subhansu Maji for helpful discussion and feedback. We would also like to thank Chengrun Yang, Xinyun Chen for their insightful comments on this manuscript. Finally, we are grateful to the following people for their contributions to creating our FRESHQA dataset: Marzena Karpinska, Dustin Tran, Daniel Cer, Sam Fullerton, Elizabeth Clark, Nishant Raj, Xiaoyu Song, Yapei Chang, Yixiao Song, Nader Akoury, Ankita Gupta, Bill Ray, Chau Pham, Wenlong Zhao, Maximilian Mozes, Simeng Sun, Ronan Salz, Kalpesh Krishna, Katherine Thai, Kanishka Misra, Salaheddin Alzu'bi, Erica Cai, Thibault Sellam, Jiao Sun, Dhruv Agarwal, Tessa Masis, Andrew Drozdov, Brian Lester, George Wei, Naveen Jafer Nizar, Shufan Wang, Youngwoo Kim, and Shib Sankar Dasgupta. This project was partially supported by award IIS-2046248 from the National Science Foundation (NSF), as well as NSF's CLOUDBANK program.

## REFERENCES

- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George Bm Van Den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego De Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, Andy Brock, Michela Paganini, Geoffrey Irving, Oriol Vinyals, Simon Osindero, Karen Simonyan, Jack Rae, Erich Elsen, and Laurent Sifre. Improving language models by retrieving from trillions of tokens. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*, volume 162 of *Proceedings of Machine Learning Research (PMLR)*, pp. 2206–2240. PMLR, 2022. URL <https://proceedings.mlr.press/v162/borgeaud22a.html>.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pp. 1877–1901, 2020. URL <https://proceedings.neurips.cc/paper/2020/file/1457c0d6bfc4967418bfb8ac142f64a-Paper.pdf>.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021a. URL <https://arxiv.org/abs/2107.03374>.
- Wenhu Chen, Xinyi Wang, William Yang Wang, and William Yang Wang. A dataset for answering time-sensitive questions. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks (NeurIPS)*, volume 1, 2021b. URL [https://datasets-benchmarks-proceedings.neurips.cc/paper\\_files/paper/2021/file/1f0e3dad99908345f7439f8ffabdfc4-Paper-round2.pdf](https://datasets-benchmarks-proceedings.neurips.cc/paper_files/paper/2021/file/1f0e3dad99908345f7439f8ffabdfc4-Paper-round2.pdf).
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. PaLM: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311*, 2022. URL <https://arxiv.org/abs/2204.02311>.
- Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Eric Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, et al. Scaling instruction-finetuned language models. *arXiv preprint arXiv:2210.11416*, 2022. URL <https://arxiv.org/abs/2210.11416>.
- Angela Fan, Yacine Jernite, Ethan Perez, David Grangier, Jason Weston, and Michael Auli. ELI5: Long form question answering. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 3558–3567, 2019. URL <https://aclanthology.org/P19-1346>.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Mingwei Chang. Retrieval augmented language model pre-training. In *Proceedings of the 37th International Conference on Machine*

- Learning (ICML)*, volume 119 of *Proceedings of Machine Learning Research (PMLR)*, pp. 3929–3938. PMLR, 2020. URL <https://proceedings.mlr.press/v119/guu20a.html>.
- Gautier Izacard, Patrick Lewis, Maria Lomeli, Lucas Hosseini, Fabio Petroni, Timo Schick, Jane Dwivedi-Yu, Armand Joulin, Sebastian Riedel, and Edouard Grave. Few-shot learning with retrieval augmented language models. 2022. URL <https://arxiv.org/abs/2208.03299>.
- Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. TriviaQA: A large scale distantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 1601–1611, 2017. URL <https://aclanthology.org/P17-1147>.
- Jungo Kasai, Keisuke Sakaguchi, Yoichi Takahashi, Ronan Le Bras, Akari Asai, Xinyan Yu, Dragomir Radev, Noah A Smith, Yejin Choi, and Kentaro Inui. Realtime qa: What’s the answer right now? 2022. URL <https://arxiv.org/abs/2207.13332>.
- Omar Khattab, Keshav Santhanam, Xiang Lisa Li, David Hall, Percy Liang, Christopher Potts, and Matei Zaharia. Demonstrate-search-predict: Composing retrieval and language models for knowledge-intensive nlp. 2022. URL <https://arxiv.org/abs/2212.14024>.
- Najoung Kim, Phu Mon Htut, Samuel R. Bowman, and Jackson Petty. (QA)<sup>2</sup>: Question answering with questionable assumptions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 8466–8487, 2023. URL <https://aclanthology.org/2023.acl-long.472>.
- Angeliki Lazaridou, Elena Gribovskaya, Wojciech Stokowiec, and Nikolai Grigorev. Internet-augmented language models through few-shot prompting for open-domain question answering. *arXiv preprint arXiv:2203.05115*, 2022.
- Brian Lester, Rami Al-Rfou, and Noah Constant. The power of scale for parameter-efficient prompt tuning. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 3045–3059, November 2021. URL <https://aclanthology.org/2021.emnlp-main.243>.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pp. 9459–9474, 2020. URL [https://proceedings.neurips.cc/paper\\_files/paper/2020/file/6b493230205f780e1bc26945df7481e5-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2020/file/6b493230205f780e1bc26945df7481e5-Paper.pdf).
- Adam Liska, Tomas Kocisky, Elena Gribovskaya, Tayfun Terzi, Eren Sezener, Devang Agrawal, Cyprien De Masson D’Autume, Tim Scholtes, Manzil Zaheer, Susannah Young, Ellen Gilsonan-Mcmahon, Sophia Austin, Phil Blunsom, and Angeliki Lazaridou. StreamingQA: A benchmark for adaptation to new knowledge over time in question answering models. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*, volume 162 of *Proceedings of Machine Learning Research (PMLR)*, pp. 13604–13622. PMLR, 2022. URL <https://proceedings.mlr.press/v162/liska22a.html>.
- Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *arXiv preprint arXiv:2307.03172*, 2023a. URL <https://arxiv.org/abs/2307.03172>.
- Nelson F Liu, Tianyi Zhang, and Percy Liang. Evaluating verifiability in generative search engines. 2023b. URL <https://arxiv.org/abs/2304.09848>.
- Shayne Longpre, Le Hou, Tu Vu, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V. Le, Barret Zoph, Jason Wei, and Adam Roberts. The flan collection: Designing data and methods for effective instruction tuning. *arXiv preprint arXiv:2301.13688*, 2023. URL <https://arxiv.org/abs/2301.13688>.
- Hongyin Luo, Yung-Sung Chuang, Yuan Gong, Tianhua Zhang, Yoon Kim, Xixin Wu, Danny Fox, Helen Meng, and James Glass. Sail: Search-augmented instruction learning. 2023. URL <https://arxiv.org/abs/2305.15225>.

- Joshua Maynez, Shashi Narayan, Bernd Bohnet, and Ryan McDonald. On faithfulness and factuality in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 1906–1919, 2020. URL <https://aclanthology.org/2020.acl-main.173>.
- Jacob Menick, Maja Trebacz, Vladimir Mikulik, John Aslanides, Francis Song, Martin Chadwick, Mia Glaese, Susannah Young, Lucy Campbell-Gillingham, Geoffrey Irving, et al. Teaching language models to support answers with verified quotes. *arXiv preprint arXiv:2203.11147*, 2022. URL <https://arxiv.org/abs/2203.11147>.
- Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, et al. Webgpt: Browser-assisted question-answering with human feedback. 2021. URL <https://arxiv.org/abs/2112.09332>.
- OpenAI. Gpt-4 technical report. *arXiv preprint arXiv:2303.0877*, 2023. URL <https://arxiv.org/abs/2303.0877>.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pp. 27730–27744, 2022. URL [https://proceedings.neurips.cc/paper\\_files/paper/2022/file/b1efde53be364a73914f58805a001731-Paper-Conference.pdf](https://proceedings.neurips.cc/paper_files/paper/2022/file/b1efde53be364a73914f58805a001731-Paper-Conference.pdf).
- Ofir Press, Muru Zhang, Sewon Min, Ludwig Schmidt, Noah A Smith, and Mike Lewis. Measuring and narrowing the compositionality gap in language models. *arXiv preprint arXiv:2210.03350*, 2022. URL <https://arxiv.org/abs/2210.03350>.
- Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research (JMLR)*, 21(140):1–67, 2020. URL <https://jmlr.org/papers/v21/20-074.html>.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761*, 2023. URL <https://arxiv.org/abs/2302.04761>.
- Zhiqing Sun, Xuezhi Wang, Yi Tay, Yiming Yang, and Denny Zhou. Recitation-augmented language models. *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*, 2023. URL <https://openreview.net/forum?id=-cqvvb-NkI>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Ed Chi, Quoc Le, and Denny Zhou. Chain of thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*, 2022. URL <https://arxiv.org/abs/2201.11903>.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. 2022. URL <https://arxiv.org/abs/2210.03629>.
- Wenhao Yu, Dan Iter, Shuohang Wang, Yichong Xu, Mingxuan Ju, Soumya Sanyal, Chenguang Zhu, Michael Zeng, and Meng Jiang. Generate rather than retrieve: Large language models are strong context generators. *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*, 2023a. URL <https://openreview.net/forum?id=fB0hRu9GZUS>.
- Wenhao Yu, Meng Jiang, Peter Clark, and Ashish Sabharwal. Ifqa: A dataset for open-domain question answering under counterfactual presuppositions. 2023b. URL <https://arxiv.org/abs/2305.14010>.
- Xinyan Yu, Sewon Min, Luke Zettlemoyer, and Hannaneh Hajishirzi. CREPE: Open-domain question answering with false presuppositions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 10457–10480, 2023c. URL <https://aclanthology.org/2023.acl-long.583>.

Michael Zhang and Eunsol Choi. SituatedQA: Incorporating extra-linguistic contexts into QA. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 7371–7387, 2021. URL <https://aclanthology.org/2021.emnlp-main.586>.

## APPENDIX

### A INTER-RATER AGREEMENT

Two authors independently evaluated a randomly sampled subset of 100 answers across models (including 50 questions with valid premises and 50 questions with false premises) in both modes RELAXED and STRICT. Table 2 reports the inter-rater agreement between the two authors in terms of exact accuracy. Overall, the high accuracies demonstrate that our evaluation protocol is reproducible and reliable.

### B EVALUATION PROTOCOL

Figure 4 shows specific examples of each evaluation criteria.

### C ADDITIONAL EXPERIMENT SETUP DETAILS FOR SECTION 3

To increase reproducibility, we select the most likely token at every decoding timestep (i.e., with a temperature of 0) and generate a maximum number of 256 tokens for all models. Note that the API for some models is non-deterministic by default, even with a temperature of 0. For non-chat models that were not pre-trained with a QA task, we feed them a text prompt of the format: “Q: <question>\nA: ” (“\n” is the new line character).

For OPENAI models, we use the 2023-03-15-preview API in AZURE OPENAI SERVICE. We use the model names text-davinci-003, code-davinci-002, gpt-3.5-turbo, and gpt-4 for GPT-3.5, CODEX, CHATGPT, and GPT-4, respectively.

### D ADDITIONAL EXPERIMENT RESULTS FOR SECTION 3

Table 3 and Table 4 show the accuracy of different LLMs on FRESHQA under STRICT (no hallucination) and RELAXED evaluations, respectively.

### E CHATGPT/GPT-4’S AWARENESS OF RECENT KNOWLEDGE

Although CHATGPT and GPT-4 were originally trained in 2021, our manual evaluation suggests that they have been exposed to data containing information beyond their knowledge cutoff of September, 2021. Figure 5 indicates that CHATGPT is aware of the recent Russian invasion of Ukraine on February 24, 2022.

### F GOOGLE SEARCH RESULTS

Figure 6 shows different types of search results from GOOGLE SEARCH for given a query.

### G A REALISTIC PROMPT FOR FRESHPROMPT

Figure 7 displays a realistic prompt for FRESHPROMPT.

### H ADDITIONAL EXPERIMENT RESULTS FOR SECTION 4

Table 5 presents the accuracy of different search engine-augmented LLMs on FRESHQA under RELAXED.



Table 2: Inter-rater agreement between two authors in terms of exact accuracy (**Acc.**) across 100 RELAXED judgements, 100 STRICT judgements, and all ALL 200 judgements. In each of these three categories, in addition to the overall accuracy (**overall**), we report accuracy across questions with valid premises (**vp**) and questions with false premises (**fp**). The high accuracies demonstrate that our evaluation protocol is reproducible and reliable.

	RELAXED			STRICT			ALL		
	overall	vp	fp	overall	vp	fp	overall	vp	fp
<b>Acc.</b>	99.0	98.0	100	96.0	100.0	92.0	97.5	99.0	96.0

ID	Question	Type	Primary answer (as of this writing)	Model response	Relaxed	Strict	Explanation
1	How old is the world's oldest verified living person?	fast-changing	116 years old	The oldest verified living person is Maria Branyas Morera, who was born on March 4, 1907, making her 116 years old.	✓	✓	the primary answer is correct, and all the provided information is accurate and up-to-date
2	When did the UK adopt the Euro?	false-premise	The UK has never adopted the Euro.	The UK has never adopted the Euro as its official currency. The country has retained the British pound sterling (GBP) as its currency.	✓	✓	the response debunks the question's false premise, and all the provided information is accurate and up-to-date
3	How much longer will Britney Griner spend in the Russian prison?	false-premise	She was released in 2022.	I'm sorry, but I have no information to suggest that Britney Griner is currently in a Russian prison.	✗	✗	a confident and definitive answer is required
4	Which language has the third largest population by number of native speakers?	slow-changing	English	1. Mandarin 2. Spanish 3. English	✓	✗	though the correct answer can be obviously inferred from the response, a response without an explicit statement is less desirable
5	Has Virginia Woolf's novel about the Ramsay family entered the public domain in the United States?	never-changing	Yes	1927's <i>To the Lighthouse</i> is in the public domain in the United States. The final answer is <i>To the Lighthouse</i> .	✗	✗	the final answer when standing alone must be correct
6	Is Tesla's stock price above \$200 a share?	fast-changing	Yes	Tesla's stock price is above \$200 a share. The stock is trading at a price of \$120.	✗	✗	conflicting information
7	When did Lionel Messi win his second FIFA World Cup?	false-premise	Messi has not won a second FIFA World Cup.	This question does not contain a valid premise. Lionel Messi has never won a FIFA World Cup.	✗	✗	despite debunking the question's false premise, the response inaccurately shapes one's knowledge perception
8	How many times has Jeremy Corbyn served as Prime Minister of the UK?	false-premise	Corbyn has never served as Prime Minister of the UK.	Zero.	✗	✗	a rebuttal is always required for false-premise questions
9	Who founded Amazon?	never-changing	Jeff Bezos	100% of people said Jeff.	✗	✗	a complete name or commonly recognized name is required
10	How many books has Colleen Hoover published?	fast-changing	25	Colleen Hoover has published over 20 books and novellas.	✗	✗	approximate numbers are not accepted unless included in the ground-truth answers
11	What is the largest model size of GPT-3?	never-changing	175B parameters	175 billion parameters. It is a 175 billion parameter, 1,377B, 137B token, 137B wordpiece, 137B BPE, 137B byte pair encoding, 137B BPE token, 137B BPE wordpiece,...	✓	✗	the primary answer is correct, but the response is ill-formed
12	How many member nations exist in the United Nations?	slow-changing	193	As of my knowledge cutoff date in September 2021, there were 193 member nations in the United Nations.	✓	✗	the primary answer is correct, but it is unclear whether the information is still up-to-date
13	Who won the biggest single-tournament payday in tennis history?	slow-changing	Novak Djokovic	The biggest single-tournament payday in tennis history was won by Novak Djokovic at the 2021 Australian Open.	✓	✗	the primary answer is correct, but the response contains hallucinated information
14	Who was the first football player to score one hundred international goals?	never-changing	Ali Daei	As of my knowledge cutoff date in September 2021, Ali Daei, a former Iranian footballer, was the first player to score 100 international goals.	✓	✗	all the information is correct and it is evident that the knowledge has not changed
15	Who is the latest winner of the Formula 1 world championship?	fast-changing	Max Verstappen	Max Verstappen won the 2022 Formula 1 world championship. He is a Dutch racing driver who competes in Formula One for Red Bull Racing. He is the son of former Formula One driver Jos Verstappen. Verstappen started his racing career in karting at the age of seven. He won the 2013 European Karting Championship and the 2013 World Karting Championship. In 2014, he moved to single-seater racing, competing in the European Formula Three Championship. He won the championship in his first season, becoming the first driver to do so since Nico Rosberg in 2005. In 2015, Verstappen moved to Formula One, driving for Toro Rosso. He became the youngest driver to compete in Formula One at the age of 17. He scored his first points in Formula One at the 2015 Hungarian Grand Prix. In 2016, Verstappen moved to Red Bull Racing. He won his first race at the 2016 Spanish Grand Prix. He became the youngest driver to win a race in Formula One at the age of 18. Verstappen finished the 2016 season in third place in the drivers' championship. In 2017, Verstappen won four races and finished the season in second place in the drivers' championship. In 2018, Verstappen won seven races and finished the season in second place in the drivers' championship. In 2019, Verstappen won nine races and finished the season in first place in the drivers' championship. He is the first Dutch driver to win the Formula One world championship.	✓	✗	the primary answer is correct, but the response contains several hallucinations (e.g., Max Verstappen did not win the Formula Three European Championship in 2014)

Figure 4: FRESHQA sample evaluation. To get credit in both evaluation modes RELAXED and STRICT, all the information in the answer must be accurate and up-to-date (examples 1 and 2). In both modes, we credit a model's response only if it provides a confident and definitive answer (example 3), or the correct answer can be obviously inferred from the response (provided all other requirements are satisfied, see example 4). The primary or final answer when standing alone must be accurate (example 5). Any additional information that is provided must not contradict the primary answer (example 6) or reshape one's perception of it (example 7). For false-premise questions, the model must point out the presence of a false premise to receive credit (example 8). For answers that involve names of entities (e.g., people), complete names or commonly recognized names are expected (example 9). Regarding numerical answers, approximate numbers are generally not accepted unless explicitly included in the ground-truth answers (example 10). In RELAXED, we accept ill-formed responses (including those in a non-English language), as well as hallucinated or outdated information that does not significantly impact the primary answer; under STRICT, however, a response that contains any hallucination, no matter how minor, will not receive credit (examples 11, 12, and 13). Furthermore, we accept a response in STRICT when the model indicates that the information might be outdated (e.g., "As of my knowledge cutoff date in September 2021") only if it is evident that the knowledge has not changed (example 14).

Table 3: Accuracy of different LLMs on FRESHQA under STRICT (no hallucination) evaluations. **Models benchmarked on the same date of April 26, 2023.** We report accuracy across different categories of questions, including *fast-changing* (*fast*), *slow-changing* (*slow*), *never-changing* (*never*), false-premise, questions that involve knowledge before 2022 ( $< 2022$ ) and since 2022 ( $\geq 2022$ ), one-hop (*1-hop*) and multi-hop (*m-hop*) questions.  $^+$  indicates a model with access to the current date.

Model (size)	knowl. cutoff	all	valid premise								false premise	
			all	fast	slow	never	< 2022	≥ 2022	1-hop	m-hop	all	< 2022
without access to a search engine												
OPENAI CODEX (N/A)	2021	25.0	<b>31.4</b>	5.6	<b>28.0</b>	60.3	<b>64.5</b>	11.5	<b>34.7</b>	23.1	5.6	7.5
GPT 3.5 (N/A)	2021	26.0	26.1	4.0	15.2	58.7	61.0	5.1	28.0	21.3	25.8	34.4
CHATGPT (N/A)	2021 <sup>+</sup>	<b>32.0</b>	28.5	7.2	16.0	61.9	63.1	7.7	29.9	25.0	<b>42.7</b>	<b>52.7</b>
GPT 4 (N/A)	2021 <sup>+</sup>	28.6	26.9	<b>12.0</b>	4.0	<b>64.3</b>	58.2	8.1	27.2	25.9	33.9	41.9
FLAN-PaLM (540B)	2022	23.4	30.3	10.4	24.8	55.6	60.3	<b>12.3</b>	32.5	25.0	2.4	3.2
PaLM (540B)	2021	7.2	9.3	0.8	11.2	15.9	20.6	2.6	9.3	9.3	0.8	1.1
w/ FEW-SHOT		20.0	26.3	5.6	19.2	54.0	56.7	8.1	25.7	<b>27.8</b>	0.8	1.1
w/ CoT		15.4	19.1	0.8	9.6	46.8	47.5	2.1	20.5	15.7	4.0	5.4
PaLMCHILLA (62B)	2022	12.2	16.0	2.4	15.2	30.2	35.5	4.3	17.2	13.0	0.8	1.1
PaLM (62B)	2021	6.2	8.2	1.6	8.8	14.3	16.3	3.4	7.8	9.3	0.0	0.0
w/ FEW-SHOT		12.8	16.8	3.2	15.2	31.7	35.5	5.5	17.9	13.9	0.8	1.1
w/ CoT		7.0	9.0	0.8	6.4	19.8	21.3	1.7	10.1	6.5	0.8	1.1
PaLM (8B)	2021	5.6	7.5	0.8	5.6	16.0	16.2	2.1	8.6	4.6	0.0	0.0
w/ FEW-SHOT		8.4	11.2	0.8	9.6	23.0	24.8	3.0	14.2	3.7	0.0	0.0
w/ CoT		7.8	10.4	0.0	6.4	24.6	24.8	1.7	11.2	8.3	0.0	0.0
FLAN-T5 XXL (11B)	2022	6.6	8.8	3.2	10.4	12.7	13.5	6.0	10.1	5.6	0.0	0.0
T5 XXL (11B)	2019	7.0	8.8	2.4	4.8	19.0	16.3	4.3	10.4	4.6	1.6	2.2
w/ FEW-SHOT		8.4	11.2	5.6	11.2	16.7	17.7	7.2	13.4	5.6	0.0	0.0
w/ CoT		6.2	8.2	2.4	6.4	15.9	15.6	3.8	8.6	7.4	0.0	0.0
T5 XL (3B)	2019	4.4	5.9	2.4	4.8	10.3	10.6	3.0	7.5	1.9	0.0	0.0
w/ FEW-SHOT		6.0	8.0	4.0	8.8	11.1	13.5	4.7	8.2	7.4	0.0	0.0
w/ CoT		2.8	3.7	2.4	1.6	7.1	7.8	1.3	4.1	2.8	0.0	0.0
T5 LARGE (770M)	2019	2.6	3.5	0.8	4.0	5.6	5.7	2.1	3.7	2.8	0.0	0.0
w/ FEW-SHOT		0.8	1.1	0.0	0.0	3.2	2.8	0.0	1.1	0.9	0.0	0.0
w/ CoT		0.8	1.1	0.8	0.0	2.4	2.1	0.4	1.1	0.9	0.0	0.0

Table 4: Accuracy of different LLMs on FRESHQA under RELAXED evaluations. **Models benchmarked on the same date of April 26, 2023.** We report accuracy across different categories of questions, including *fast-changing* (*fast*), *slow-changing* (*slow*), *never-changing* (*never*), false-premise, questions that involve knowledge before 2022 ( $< 2022$ ) and since 2022 ( $\geq 2022$ ), one-hop (*1-hop*) and multi-hop (*m-hop*) questions. <sup>+</sup> indicates a model with access to the current date.

Model (size)	knowl. cutoff	all	valid premise								false premise	
			all	fast	slow	never	< 2022	≥ 2022	1-hop	m-hop	all	< 2022
without access to a search engine												
OPENAI CODEX (N/A)	2021	25.6	32.2	6.4	29.6	60.3	66.0	11.9	35.4	24.1	5.6	7.5
GPT 3.5 (N/A)	2021	32.4	32.4	8.0	28.0	61.1	68.1	11.1	34.7	26.9	32.3	43.0
CHATGPT (N/A)	2021 <sup>+</sup>	41.4	36.7	10.4	32.8	66.7	76.6	12.8	36.2	38.0	55.6	66.7
GPT 4 (N/A)	2021 <sup>+</sup>	<b>46.4</b>	<b>39.6</b>	<b>14.4</b>	<b>35.2</b>	<b>69.0</b>	<b>80.9</b>	<b>14.9</b>	<b>39.2</b>	<b>40.7</b>	<b>66.9</b>	<b>83.9</b>
FLAN-PALM (540B)	2022	23.6	30.3	10.4	24.8	55.6	60.3	12.3	32.5	25.0	3.2	4.3
PALM (540B)	2021	12.2	16.0	2.4	14.4	31.0	34.8	4.7	16.4	14.8	0.8	1.1
w/ FEW-SHOT		20.2	26.3	5.6	19.2	54.0	56.7	8.1	25.7	27.8	1.6	2.2
w/ CoT		22.8	28.2	4.0	20.0	60.3	64.5	6.4	28.4	27.8	6.5	8.6
PALMCHILLA (62B)	2022	15.0	19.4	2.4	19.2	36.5	43.3	5.1	20.1	17.6	1.6	2.2
PALM (62B)	2021	8.6	11.2	2.4	11.2	19.8	22.0	4.7	11.6	10.2	0.8	1.1
w/ FEW-SHOT		14.2	18.4	4.0	15.2	35.7	39.0	6.0	18.7	17.6	1.6	2.2
w/ CoT		12.8	16.2	2.4	15.2	31.0	34.8	5.1	17.5	13.0	2.4	3.2
PALM (8B)	2021	8.8	11.2	0.8	11.2	21.6	21.1	5.2	13.1	6.5	1.6	2.1
w/ FEW-SHOT		9.2	12.2	0.8	10.4	25.4	27.0	3.4	15.3	4.6	0.0	0.0
w/ CoT		11.4	15.2	2.4	11.2	31.7	32.6	4.7	16.8	11.1	0.0	0.0
FLAN-T5 XXL (11B)	2022	7.2	9.6	3.2	12.0	13.5	14.2	6.8	10.8	6.5	0.0	0.0
T5 XXL (11B)	2019	10.8	13.8	3.2	12.8	25.4	22.7	8.5	16.0	8.3	1.6	2.2
w/ FEW-SHOT		9.0	12.0	5.6	11.2	19.0	19.1	7.7	14.6	5.6	0.0	0.0
w/ CoT		13.0	17.3	4.0	17.6	30.2	31.2	8.9	19.0	13.0	0.0	0.0
T5 XL (3B)	2019	5.8	7.7	4.0	5.6	13.5	13.5	4.3	9.0	4.6	0.0	0.0
w/ FEW-SHOT		6.0	8.0	4.0	8.8	11.1	13.5	4.7	8.2	7.4	0.0	0.0
w/ CoT		5.2	6.9	3.2	4.0	13.5	14.2	2.6	8.6	2.8	0.0	0.0
T5 LARGE (770M)	2019	4.4	5.3	2.4	4.8	8.7	7.1	4.3	5.6	4.6	1.6	2.2
w/ FEW-SHOT		0.8	1.1	0.0	0.0	3.2	2.8	0.0	1.1	0.9	0.0	0.0
w/ CoT		2.2	2.9	0.8	0.8	7.1	7.1	0.4	3.4	1.9	0.0	0.0



Tell me about the recent Russian invasion of Ukraine.



As an AI language model, I can provide an objective summary of the recent events that have taken place between Russia and Ukraine.

In late February 2022, Russia began a military invasion of Ukraine, citing security concerns and allegations of Ukrainian aggression towards Russian-speaking populations in eastern Ukraine. The invasion quickly escalated into a full-scale war, with Russian forces pushing deep into Ukrainian territory and capturing several key cities.

The conflict has been marked by intense fighting, heavy casualties on both sides, and numerous reports of human rights abuses and war crimes. The international community has condemned Russia's actions and imposed economic sanctions in response, but efforts to negotiate a ceasefire or peaceful resolution to the conflict have so far been unsuccessful.

As of April 2023, the situation remains tense and the conflict ongoing, with no clear end in sight. The toll on civilian populations in both Russia and Ukraine has been significant, and the humanitarian situation in many parts of Ukraine is dire.



When did Russia invade Ukraine recently?



Russia invaded Ukraine on February 24, 2022.

Figure 5: CHATGPT is aware of the recent Russian invasion of Ukraine on February 24, 2022. Questions asked on April 9, 2023.

The image shows two side-by-side Google search result pages. The left page is for the query "Who is the most recent American Idol winner?". It features an "answer box" with a photo of Noah Thompson and text stating he is the latest winner. Below this are "related questions" such as "Who is the newest American Idol winner?" and "Who won American Idol 2023 last night?". At the bottom are "organic results" from GoldDerby.com listing American Idol winners by season. The right page is for the query "What is the name of the first animal to land on the moon?". It features an "answer box" stating that steppe tortoises were the first animals to reach the moon. Below this are "related questions" like "What was the first animal to survive in space?". Further down are "questions and answers" from crowdsourced platforms like Study.com and Quora. At the bottom are "organic results" from the Royal Museums Greenwich and Homework.Study.com, both mentioning the dog Laika as the first animal in space.

Figure 6: GOOGLE SEARCH produces different types of search results for given a query, including the *answer box*, *organic results*, and other useful information, such as the *knowledge graph*, *questions and answers* from crowdsourced QA platforms, and *related questions* that search users also ask. Each of these results contains an associated *text snippet* along with other information, such as *source webpage*, *date*, *title*, and *highlighted words*.





Figure 7: A realistic prompt for FRESHPROMPT. We cast all retrieved evidences into a unified format with useful information, including source webpage, date, title, text snippet, and highlighted words. Few-shot demonstrations are provided at the beginning of the prompt. Each demonstration shows the model an example question and a list of retrieved evidences for the question, followed by some reasoning over the evidences to figure out the most relevant and up-to-date answer.

Table 5: Accuracy of different search engine-augmented LLMs on FRESHQA under RELAXED evaluations. **Models benchmarked on the same date of April 26, 2023.** We report accuracy across different categories of questions, including *fast-changing* (*fast*), *slow-changing* (*slow*), *never-changing* (*never*), false-premise, questions that involve knowledge before 2022 ( $< 2022$ ) and since 2022 ( $\geq 2022$ ), one-hop (*1-hop*) and multi-hop (*m-hop*) questions. <sup>+</sup> indicates a model with access to the current date. UTD stands for “up-to-date”.

Model	knowl. cutoff	all	valid premise								false premise	
			all	fast	slow	never	< 2022	≥ 2022	1-hop	m-hop	all	< 2022
<i>comparison against baselines</i>												
GOOGLE SEARCH	UTD	47.4	58.8	42.4	56.0	77.8	74.5	49.4	66.4	39.8	12.9	11.8
GPT-3.5	2021	32.4	32.4	8.0	28.0	61.1	68.1	11.1	34.7	26.9	32.3	43.0
GPT-3.5 + SELF-ASK	UTD	42.0	51.6	36.8	44.8	73.0	74.5	37.9	53.0	48.1	12.9	17.2
GPT-3.5 + FRESHPROMPT	UTD	62.0	68.9	51.2	70.4	84.9	78.0	63.4	75.0	53.7	41.1	49.5
PPLX.AI	UTD	66.2	68.9	48.8	67.2	90.5	85.1	59.1	76.1	50.9	58.1	60.2
GPT-4	2021 <sup>+</sup>	46.4	39.6	14.4	35.2	69.0	80.9	14.9	39.2	40.7	66.9	83.9
GPT-4 + SELF-ASK	UTD	50.4	48.4	40.0	49.6	55.6	52.5	46.0	45.1	56.5	56.5	69.9
GPT-4 + FRESHPROMPT	UTD	<b>77.8</b>	<b>78.7</b>	<b>61.6</b>	<b>79.2</b>	<b>95.2</b>	<b>90.8</b>	<b>71.5</b>	<b>83.2</b>	<b>67.6</b>	<b>75.0</b>	<b>80.6</b>
<i>sensitivity and ablation studies</i>												
GPT-3.5	2021	32.4	32.4	8.0	28.0	61.1	68.1	11.1	34.7	26.9	32.3	43.0
GPT-3.5 + FRESHPROMPT	UTD	62.0	68.9	51.2	70.4	84.9	78.0	63.4	75.0	53.7	41.1	49.5
w/ PREMISE CHECK	UTD	41.0	33.5	23.2	32.0	45.2	44.0	27.2	37.7	23.1	63.7	72.0
GPT-4	2021 <sup>+</sup>	46.4	39.6	14.4	35.2	69.0	80.9	14.9	39.2	40.7	66.9	83.9
GPT-4 w/ SNIPPETS ONLY & SEARCH ORDER	UTD	77.6	78.2	59.2	<b>80.0</b>	95.2	90.8	70.6	82.1	68.5	75.8	83.9
GPT-4 w/ SNIPPETS ONLY & TIME ORDER	UTD	77.6	78.2	59.2	79.2	<b>96.0</b>	90.1	71.1	82.1	68.5	75.8	86.0
GPT-4 w/ SNIPPETS ONLY & RANDOM ORDER	UTD	75.4	76.1	58.4	73.6	<b>96.0</b>	90.8	67.2	80.6	64.8	73.4	81.7
GPT-4 + FRESHPROMPT	UTD	77.8	78.7	61.6	79.2	95.2	90.8	71.5	<b>83.2</b>	67.6	75.0	80.6
w/ PREMISE CHECK	UTD	78.8	76.3	59.2	76.8	92.9	87.2	69.8	82.1	62.0	<b>86.3</b>	<b>90.3</b>
w/o ANSWER BOX	UTD	76.2	76.6	59.2	76.0	94.4	90.1	68.5	81.0	65.7	75.0	80.6
w/o ANSWER BOX & RELEVANT INFO	UTD	74.8	75.0	56.0	74.4	94.4	89.4	66.4	80.6	61.1	74.2	81.7
w/ 1 EVIDENCE	UTD	67.2	67.3	47.2	66.4	88.1	85.8	56.2	72.0	55.6	66.9	79.6
w/ 5 EVIDENCES	UTD	74.2	75.0	56.8	74.4	93.7	87.2	67.7	81.7	58.3	71.8	77.4
w/ 15 EVIDENCES	UTD	<b>79.0</b>	<b>79.5</b>	<b>62.4</b>	<b>80.0</b>	<b>96.0</b>	90.1	<b>73.2</b>	<b>83.2</b>	<b>70.4</b>	77.4	81.7
w/ 15 DEMONSTRATIONS	UTD	77.2	78.2	60.0	78.4	<b>96.0</b>	<b>91.5</b>	70.2	82.8	66.7	74.2	79.6
w/ LONG DEMONSTRATION ANSWERS	UTD	77.8	77.9	60.8	77.6	95.2	90.1	70.6	82.8	65.7	77.4	83.9